Questions

1. **(Basic Knowledge) Tuning the Regularization Parameter**
   - What does the regularization parameter λ control in lasso/ridge/elastic net?
   - How is λ typically selected in practice?
   - How does changing λ affect model complexity and the bias–variance tradeoff?
2. **(Conceptual Thinking) Prediction vs. Interpretation**
   - Why might the λ that gives the best prediction not be ideal for interpretation or variable selection?
   - What is the difference between choosing λ for **prediction accuracy** versus **variable selection**?
3. **(Technical Understanding) Statistical Inference after Regularization**
   - Why is classical statistical inference (e.g., standard errors, p-values) not directly valid after fitting a lasso model?
4. (**Open Question**) In your view, when is regularized regression best treated as a **predictive tool** rather than an **inferential model**?